

# Proceedings Track

## Semantic-Aware Logit Adjustment for Long-Tailed Recognition

**Editors:** List of editors' names

### Abstract

Training deep neural networks on imbalanced data is difficult as conventional training methods exhibit bias towards the head classes. Existing methods address long-tailed bias issue through frequency correction or representation learning, but they do not quantify class-level feature overlap and optimize the loss accordingly. In this work, we introduce Semantic-Aware Logit Adjustment (SALA), a training loss which integrates the concept of frequency priors together with semantic crowdedness penalty based on the similarities of the classifiers' prototypes. Our method not only considers the rarity of each class, but also the degree of overlap of the representations between the classes to optimize the most structurally crowded classes. Extensive experiments performed on standard benchmarks like ImageNet-LT, CIFAR-10-LT, CIFAR-100-LT and iNaturalist 2018 show highly competitive performance, particularly excelling in rare-class accuracy.

### 1. Introduction

Deep learning models have been very successful in dealing with well-curated datasets with even class distributions (Schmidhuber, 2015). However, in practice, such as in applications like self-driving cars (Liu et al., 2024), medical diagnoses (Jiménez-Sánchez et al., 2025), and social media analysis (Turcan and McKeown, 2019) often exhibit an inherent class imbalance, where a handful of head classes comprise the majority of the data, while many tail classes lack representation. Deep learning models trained using regular cross-entropy often experience severe performance drops for the tail classes.

To address this performance deterioration, the initial set of methods concentrate on resampling the dataset (He et al., 2008; Chawla et al., 2002; Guo and Viktor, 2004; Douzas and Bação, 2017) and cost-sensitive re-weighting (Lin et al., 2018; Cui et al., 2019; Park et al., 2021; Yu et al., 2022). While resampling attempts to equalize class distributions through over-sampling of the tail or under-sampling of the head class, it often results in a substantial risk of overfitting to the minority class or wasteful utilization of the abundant majority class samples. In contrast, re-weighting approaches increase penalties for tail class errors, which could lead to optimization instabilities and cause degradation in accuracy for the majority classes. The critical problem with this type of approach is that it tries to influence the training process without solving the bias issue that arises from the imbalance itself. More recent methods focus on modifying the class margins (Cho and Krähenbühl, 2023; Kato and Hotta, 2023; Deng et al., 2021; Menon et al., 2021) or reshaping class geometry (Tao et al., 2023; Zhou et al., 2025; Chen et al., 2025; Li et al., 2024; Zhang et al., 2024a), while some rely on contrastive objectives (Li et al., 2022b; Zhu et al., 2022; Du et al., 2024; Hou et al., 2023) to improve inter-class separation. However, these approaches largely treat class imbalance and representation separation as broadly shared across the classes without explicitly accounting for the varying degree of feature overlap between the classes.

# Proceedings Track

In practice, semantically similar rare classes require greater attention so this motivates a more adaptive mechanism that accounts for both class rarity and semantic crowdedness.

To address these concerns, we introduce Semantic-Aware Logit Adjustment which accounts for both class rarity and semantic crowdedness. This design provides our framework two perspectives. First, it preserves the benefits of frequency-aware correction while avoiding the limitation of treating all classes as equally difficult. Second, it introduces an explicit mechanism for addressing the residual semantic bias that persists even after frequency adjustment. By penalizing the crowded classes more strongly, the proposed approach encourages the feature extractor to learn more discriminative representations for semantically similar classes. As a result, the learned features become more semantically organized, enabling minority classes gain clearer separation from visually similar neighbors while maintaining stable optimization by grounding the semantic penalty in the geometry of the current classifier rather than in a separate auxiliary objective.

The major contributions are:

- To propose Semantic-Aware Logit Adjustment (SALA), a new loss that incorporates both the class rarity and semantic crowdedness. By considering both aspects of the problem, we are able to prevent the excessive correction of semantically distinct classes while simultaneously optimizing the structurally crowded classes.
- SALA is a lightweight, plug-and-play loss-level modification, seamlessly integrated into existing long-tailed recognition pipelines without requiring additional networks, contrastive branches, or complex training architectures.
- To perform experiments on benchmark datasets, such as ImageNet-LT ([Russakovsky et al., 2015](#)), iNaturalist 2018 ([Horn et al., 2018](#)), CIFAR-10-LT, CIFAR-100-LT ([Krizhevsky, 2009](#)). Our proposed method demonstrates its effectiveness by achieving competitive overall performance and substantial improvements specifically for rare classes.

## 2. Related Works

### 2.1. Re-sampling and Re-weighting Methods

Traditional approaches for fixing class imbalance either modify how the training data is distributed or adjust loss function weights. To mitigate class imbalance, re-sampling techniques can be employed to over-sample minority classes or under-sample majority classes ([Chawla et al., 2002](#); [Mujahid et al., 2024](#); [He et al., 2008](#); [Guo and Viktor, 2004](#); [Douzas and Baao, 2017](#)). However, over-sampling minority class can lead to overfitting, while under-sampling majority class results in the loss of potentially important information. Alternatively, re-weighting assigns greater importance to the minority classes by adjusting their weight during loss calculation ([Lin et al., 2018](#); [Cui et al., 2019](#); [Park et al., 2021](#); [Yu et al., 2022](#)). This method preserves the cardinality of the dataset but it can destabilize optimization, particularly under severe class imbalance where gradients from rare classes may become disproportionately dominant.

# Proceedings Track

## 2.2. Logit Adjustment and Margin-Based Approaches

Margin-based methods (Cho and Krähenbühl, 2023; Kato and Hotta, 2023; Deng et al., 2021) tweak the decision boundary rather than the training data itself. For example, Equalization Loss (Tan et al., 2020), Seesaw Loss (Wang et al., 2021), and LDAM (Cao et al., 2019) use class-specific margins to create broader decision boundaries, especially for minority classes. Logit Adjustment (LA) (Menon et al., 2021) and Balanced Softmax (Ren et al., 2020) then offer a solid Bayesian backing for this approach, adjusting the logits by  $\log \pi_c$  to directly address the shifts in class priors. More recently, researchers enhanced the LA framework with extra awareness terms. Take Adaptive Logit Adjustment (ALA) (Zhao et al., 2021), which adjusts the frequency term based on instance difficulty. Then there’s Difficulty-aware Balancing Margin (DBM) Loss (Son et al., 2024), which modifies the margin penalty to handle different learning difficulties across classes. Likewise, Margin Calibration (MARC) (Wang et al., 2022) works on logits through affine transformations. Many recent works (Tao et al., 2023; Zhou et al., 2025; Chen et al., 2025; Li et al., 2024; Zhang et al., 2024a) see difficulty as variations or errors in statistics at the instance level. In contrast, SALA is different since it goes beyond these by assessing the geometric entanglement.

## 2.3. Contrastive Learning for Long-Tailed Recognition

Supervised Contrastive Learning (SupCon) (Khosla et al., 2021) improves representation learning by pulling samples from the same class together while separating other classes. Under long-tailed distributions, head classes dominate contrastive interactions and bias the learned representations. Few works introduce the class- or structure-aware contrastive objectives, including Balanced Contrastive Learning (BCL) (Zhu et al., 2022), Targeted Supervised Contrastive Learning (TSC) (Li et al., 2022b), and Subclass-balancing Contrastive Learning (SBCL) (Hou et al., 2023). More recent methods further address this bias through decoupled objectives (Qiu et al., 2024), probabilistic modeling of class distributions (Du et al., 2024). Despite aforementioned advances, existing methods mainly modify sample-level contrastive interactions or their optimization. But SALA explicitly measures the semantic crowdedness of class prototypes and incorporates this class-specific geometry into logit adjustment which is just a plug-and-play modification.

## 3. Methodology

### 3.1. Preliminary

Let  $\mathcal{D} = \{(x_i, y_i)\}_{i=1}^N$  be a long-tailed dataset where  $y_i \in [C] := \{1, \dots, C\}$ . The sample counts  $n_c$  follow  $n_1 \geq \dots \geq n_C > 0$  and the empirical class prior is  $\pi_c = n_c/N$ . A network maps inputs to features  $h = f_\theta(x) \in \mathbb{R}^d$  and logits  $z = Wh \in \mathbb{R}^C$ , where  $W$  contains prototypes  $w_c^\top$ . We denote  $L_2$ -normalized vectors as  $\bar{w}_c = w_c/\|w_c\|_2$ .

The standard cross-entropy (CE) loss is expressed by:

$$\mathcal{L}_{\text{CE}}(z, y) = -\log \frac{\exp(z_y)}{\sum_{j=1}^C \exp(z_j)} \quad (1)$$

Standard CE optimization works great with balanced datasets, but it tends to model the skewed training prior  $\mathbb{P}(y) \approx \pi_y$ . This becomes a concern in long-tailed distributions

# Proceedings Track

where it favors head classes, leaving tail classes behind. Menon et al. (2021) came up with Logit Adjustment (LA) technique to fix this issue. Inspired by statistical decision theory, LA tweaks the logits based on class priors to even out the error rates and give greater importance to the rare classes. Their method, which focuses on increasing the margin for less common classes, is summarized by the LA loss formula:

$$\mathcal{L}_{\text{LA}}(z, y) = -\log \frac{\exp(z_y + \tau \log \pi_y)}{\sum_{j=1}^C \exp(z_j + \tau \log \pi_j)} \quad (2)$$

where  $\tau > 0$  is a scaling parameter (often set to  $\tau = 1$ ).

To understand the exact mechanics of this adjustment on head versus tail classes, by considering gradient of  $\mathcal{L}_{\text{LA}}$  w.r.t. the target logit  $z_y$ . Let  $p_y^{\text{LA}}$  show the adjusted softmax probability for the target class derived from the Equation 3.1. The gradient is given by:

$$\frac{\partial \mathcal{L}_{\text{LA}}}{\partial z_y} = p_y^{\text{LA}} - 1 \quad (3)$$

Let us consider two training samples, one from head class and another from tail class, they both yield the exact same raw target logit  $z$ . Due to the frequency-based priors,

$$z + \tau \log \pi_{y_t} < z + \tau \log \pi_{y_h} \quad (4)$$

Consequently, the adjusted probability for the tail class is strictly suppressed compared to head class ( $p_{y_t}^{\text{LA}} < p_{y_h}^{\text{LA}}$ ). Because the magnitude of the gradient updating the logit is  $|p_y^{\text{LA}} - 1| = 1 - p_y^{\text{LA}}$ , the tail sample produces a larger gradient scale:

$$\left| \frac{\partial \mathcal{L}_{\text{LA}}}{\partial z_{y_t}} \right| > \left| \frac{\partial \mathcal{L}_{\text{LA}}}{\partial z_{y_h}} \right| \quad (5)$$

This enables the model to adjust its focus dynamically, prioritizing the optimization for harder, underrepresented classes. It pushes for bigger gradient steps where they are needed the most.

We analyze the geometry of the learned classifier through the Neural Collapse (NC) framework (Papayan et al., 2020). In terminal phase of training, NC1 suggests that within-class features become tightly concentrated around their class means, while NC3 indicates an alignment between the classifier weights and the corresponding class means. Accordingly, for a sample  $(x, y)$ , its feature representation can be approximated as

$$h(x) \approx s \cdot \bar{w}_y, \quad (6)$$

where  $s > 0$  denotes the feature norm and  $\bar{w}_y$  is the normalized prototype of the target class. Under this approximation, the logit corresponding to class  $c$  becomes

$$z_c(x) = w_c^\top h(x) \approx s \|w_c\|_2 \left( \bar{w}_c^\top \bar{w}_y \right). \quad (7)$$

Thus, the relative logits are directly determined by the angular relationships between the class prototypes. In an ideally separated representation, non-target prototypes have low cosine similarity with the target prototype. But under long-tailed training, this geometry can become distorted, mainly for the minority classes whose prototypes may become

# Proceedings Track

closely aligned with the dominant classes. Such positive prototype overlap increases the competition faced by the corresponding class and leaves a residual geometric bias that frequency-based adjustment alone cannot capture.

## 3.2. Semantic-Aware Logit Adjustment (SALA)

We introduce Semantic-Aware Logit Adjustment (SALA) which penalizes classes based on feature overlap in latent space rather than training data rarity alone. This allows the model to determine how difficult it is to distinguish between classes during learning independently of how frequently each class appear.

The Semantic Crowdedness Score  $\Omega_c$  measures how tightly packed and intertwined a specific class is on the unit hypersphere. By the Neural Collapse approximation of Eq. (6), the normalized prototypes  $\bar{w}_c = w_c / \|w_c\|_2$  from Section 3.1 represent the class-mean directions in feature space so crowdedness measured on prototypes is crowdedness measured on the underlying class clusters. Thus, we quantify the alignment between any two classes  $c$  and  $k$  with the same cosine similarity term that drove the logit competition in Eq. (7),  $\bar{w}_c^\top \bar{w}_k = \cos(\theta_{c,k})$ , where  $\theta_{c,k}$  is the angle between the two prototypes. We define the semantic crowdedness score as:

$$\Omega_c = \sum_{k \neq c} \max(0, \bar{w}_c^\top \bar{w}_k) \quad (8)$$

This formulation introduces two main mathematical design rules. First, the rectified linear unit  $\max(0, \cdot)$  retains only the active geometric interference. Classes separated by an angle of  $90^\circ$  or more do not affect each other since their similarity is zero. But for angles under  $90^\circ$ , the impact grows with how close the classes are. Second,  $\Omega_c$  deliberately excludes the empirical frequency prior  $\pi_k = n_k/N$  that drives  $\mathcal{L}_{LA}$  in Eq. (3.1). It serves as a pure measure of how similar the classes are, orthogonal to the frequency term it sits alongside in Eq. (9) below. This makes it extra sensitive to those rare class entanglements where sample counts almost match up ( $\pi_{y_t}/\pi_k \rightarrow 1$ ).

We integrate this semantic term directly into the logit adjustment framework of Eq. (3.1) to penalize the crowded classes during optimization. The final SALA loss simply augments  $\mathcal{L}_{LA}$  with the  $-\alpha\Omega_y$  term:

$$\mathcal{L}_{SALA}(z, y) = -\log \frac{\exp(z_y + \tau \log \pi_y - \alpha\Omega_y)}{\sum_{j=1}^C \exp(z_j + \tau \log \pi_j - \alpha\Omega_j)} \quad (9)$$

Here  $\tau > 0$  retains its role from Eq. (3.1), controlling the frequency prior adjustment, while  $\alpha > 0$  is a new scaling hyperparameter for the semantic penalty. Subtracting  $\alpha\Omega_y$  from the target logit reduces the model’s confidence for crowded classes. These changes encourage the feature extractor to move overlapping class representations away from each other.

## 4. Experiments

### 4.1. Datasets

The proposed SALA loss is assessed on four commonly employed benchmarks for the task of long-tailed visual recognition, which include CIFAR-10-LT (Krizhevsky, 2009), CIFAR-100-LT (Krizhevsky, 2009), ImageNet-LT (Liu et al., 2019), and iNaturalist2018 (Horn et al.,




Figure 1: Split-wise testing accuracy over training iterations on CIFAR-100-LT for imbalance ratios 50, 100, and 200, respectively, where the blue line denotes the method with the semantic crowdedness term and the orange line denotes the variant with just the frequency term.

2018). The benchmark datasets involve long-tailed synthetic and real-world data at both small-scale and large-scale levels, providing an extensive experiment under different levels of imbalance severity. For the CIFAR-10-LT and CIFAR-100-LT datasets, long-tailed versions of the datasets are obtained by exponential undersampling their counterparts. Three levels of imbalance factor ( $IF$ ) are considered, that are  $\{50, 100, 200\}$  where  $IF$  represents the ratio of the most common class to the least common class. On the other hand, ImageNet-LT and iNaturalist2018 are large-scale naturally imbalanced datasets exhibiting highly skewed class distributions. Dataset statistics are summarized in Table 2.

# Proceedings Track

Table 1: Top-1 accuracy (%) on CIFAR-10-LT and CIFAR-100-LT using ResNet32.

| Dataset<br>Imbalance Factor                | CIFAR-10-LT  |              |              | CIFAR-100-LT |              |              |
|--------------------------------------------|--------------|--------------|--------------|--------------|--------------|--------------|
|                                            | 50           | 100          | 200          | 50           | 100          | 200          |
| Cross-Entropy                              | –            | 77.40        | 71.20        | 43.90        | 38.43        | 34.84        |
| Logit Adjustment (Menon et al., 2021)      | 83.98        | 81.03        | 77.97        | 52.11        | 46.90        | 43.36        |
| LDAM-DRW (Cao et al., 2019)                | 81.03        | 77.03        | 73.52        | 47.62        | 42.04        | 38.91        |
| Balanced Softmax (Ren et al., 2020)        | –            | 82.90        | 75.50        | 50.50        | 45.00        | 43.30        |
| cRT (Kang et al., 2020)                    | 84.21        | 79.15        | 73.06        | 50.86        | 45.12        | 41.73        |
| Weight Balancing (Alshammari et al., 2022) | –            | –            | –            | 49.80        | 44.30        | 40.80        |
| MisLAS (Zhong et al., 2021)                | 85.70        | 82.10        | –            | 52.30        | 47.00        | 42.33        |
| GCL (Li et al., 2022a)                     | 85.46        | 82.68        | 79.03        | 53.55        | 48.71        | 44.88        |
| MetaSAug (Li et al., 2021)                 | 84.03        | 80.54        | 76.89        | 51.90        | 46.87        | 39.94        |
| SWRM (Chen et al., 2024)                   | –            | –            | –            | 52.01        | 46.80        | –            |
| LADE (Hong et al., 2021)                   | –            | –            | –            | 50.50        | 45.40        | –            |
| ResLT (Jiequan et al., 2022)               | –            | 83.46        | <b>80.44</b> | –            | 49.98        | 45.34        |
| BKD (Zhang et al., 2023)                   | 85.10        | 82.50        | –            | 51.70        | 46.50        | –            |
| TSC (Li et al., 2022b)                     | 82.90        | 79.70        | –            | 47.40        | 43.80        | –            |
| SETF (Yang et al., 2022)                   | 83.80        | 79.70        | 74.20        | 52.50        | 48.30        | 43.00        |
| ARBloss (Xie et al., 2023)                 | <b>86.10</b> | 83.80        | 79.50        | 49.70        | 47.10        | 42.70        |
| CAL (Yan et al., 2024)                     | 85.90        | <b>84.10</b> | 80.00        | <b>54.30</b> | <b>50.10</b> | <b>46.50</b> |
| <b>SALA</b>                                | <b>86.25</b> | <b>84.96</b> | <b>81.11</b> | <b>55.19</b> | <b>51.41</b> | <b>47.30</b> |

Table 2: Datasets Statistics.

| Dataset         | Classes | Train   | Test   |
|-----------------|---------|---------|--------|
| CIFAR-10        | 10      | 50,000  | 10,000 |
| CIFAR-100       | 100     | 50,000  | 10,000 |
| ImageNet-LT     | 1,000   | 115,846 | 50,000 |
| iNaturalist2018 | 8,142   | 437,513 | 24,426 |

## 4.2. Implementation Details

The same training protocol from Ren et al. (2020) is followed for all experiments. We use Nesterov SGD with a momentum of 0.9 and weight decay of  $5 \times 10^{-4}$  for all model weights. For CIFAR-10-LT and CIFAR-100-LT, we use ResNet-32 as backbone, with a mini-batch size of 512. The learning rate starts at 0.05, increasing linearly to 0.1 over the first 800 iterations and then follows a cosine annealing schedule for 200 epochs, eventually dropping to 0. For iNaturalist 2018 and ImageNet-LT, we switch to ResNet-50 and train from scratch for 90 epochs using the same optimizer setup. Experiments were conducted on the NVIDIA A100 GPU.

## 4.3. Main Results

**CIFAR-10-LT.** Regarding CIFAR-10-LT, SALA outperforms other existing approaches across all imbalance factors. When  $IF = 50$ , SALA achieves an accuracy of 86.25% by

# Proceedings Track

Table 3: Top-1 accuracy (%) on ImageNet-LT and iNat2018 using ResNet50.

| Dataset                               | ImageNet-LT  | iNat2018     |
|---------------------------------------|--------------|--------------|
| Cross-Entropy                         | 44.51        | 63.80        |
| Logit Adjustment (Menon et al., 2021) | 52.75        | 68.11        |
| LDAM-DRW (Cao et al., 2019)           | 48.80        | 68.00        |
| Balanced Softmax (Ren et al., 2020)   | 47.20        | –            |
| cRT (Kang et al., 2020)               | 50.30        | 65.20        |
| MisLAS (Zhong et al., 2021)           | 53.40        | 71.60        |
| Seesaw Loss (Wang et al., 2021)       | 50.40        | –            |
| MetaSAug (Li et al., 2021)            | 47.39        | 68.75        |
| LADE (Hong et al., 2021)              | 53.00        | 70.00        |
| GCL (Li et al., 2022a)                | <b>54.88</b> | <b>72.01</b> |
| GALA (Zhang et al., 2024b)            | 53.40        | 71.20        |
| ALA (Zhao et al., 2021)               | 52.40        | 70.70        |
| MARC (Wang et al., 2022)              | 52.30        | 70.40        |
| SETF (Yang et al., 2022)              | 44.70        | –            |
| TSC (Li et al., 2022b)                | 52.40        | 69.70        |
| ARBloss (Xie et al., 2023)            | 52.80        | –            |
| CAL (Yan et al., 2024)                | 49.70        | –            |
| <b>SALA</b>                           | <b>56.02</b> | <b>74.34</b> |

outperforming other robust alternatives like ARBloss (86.10%), MisLAS (85.70%), and GCL (85.46%). As the degree of imbalance increases, SALA consistently maintains the lead, by producing 84.96% and 81.11% accuracy for  $IF$  values of 100 and 200, respectively. These accuracies significantly outperform all the other competing techniques, where CAL reports accuracies of 84.10% and 80.00% for  $IF$  values of 100 and 200, respectively.

**CIFAR-100-LT.** SALA demonstrates even more significant improvements on the CIFAR-100-LT dataset. As we can see in Table 1 and from the split-wise learning curves in Fig. 1, SALA activation quickly leads to overall balanced improvements in performance. The proposed approach outperforms current approaches like CAL (54.30%), GCL (53.55%), MisLAS (52.30%), and SWRM (52.01%) for  $IF = 50$ . As imbalance factor increases, SALA continues to outperform its competitors. Specifically, SALA achieves accuracies of 51.41% at  $IF = 100$  and 47.30% at  $IF = 200$ , both being competitive to other existing methods.

**ImageNet-LT.** We further investigate the performance of our method by conducting experiments on ImageNet-LT (Table 3) to examine the efficacy of our approach under the setting of extreme long-tailed class distributions. It is shown that SALA can achieve a Top-1 accuracy of 56.02%, which is significantly better than other strong baselines, that includes GCL (54.88%), MisLAS (53.40%), LADE (53.00%), and GALA (53.40%).

**iNaturalist2018.** We further evaluate our method on the iNaturalist2018 dataset. As we can see in Table 3 that our method obtains 74.34% Top-1 accuracy and outperforms other existing approaches like GCL (72.01%), MisLAS (71.60%) and GALA (71.20%). These

# Proceedings Track

improvements exhibit that our method can effectively address the overconfidence problem of tail classes while learning discriminative features for the fine-grained categories.

Table 4: Ablation study on CIFAR-100-LT (ResNet32).  $\checkmark$  indicates that the term is active. Top-1 accuracy (%) reported.

| Configuration      | Freq. Term<br>( $\tau \log \pi_c$ ) | Geom. Term<br>( $\alpha \Omega_c$ ) | CIFAR-100-LT |              |              |
|--------------------|-------------------------------------|-------------------------------------|--------------|--------------|--------------|
|                    |                                     |                                     | IF=50        | IF=100       | IF=200       |
| Baseline           | $\times$                            | $\times$                            | 43.90        | 38.43        | 34.84        |
| Freq. Only         | $\checkmark$                        | $\times$                            | 52.11        | 46.90        | 43.36        |
| Sem. Only          | $\times$                            | $\checkmark$                        | 49.73        | 44.85        | 41.57        |
| <b>SALA (Both)</b> | $\checkmark$                        | $\checkmark$                        | <b>55.19</b> | <b>51.41</b> | <b>47.30</b> |

## 4.4. Ablation Study

To isolate the contribution of each component in SALA, we evaluate four configurations on CIFAR-100-LT: (1) standard cross-entropy, (2) frequency prior only, (3) semantic crowdedness penalty only, and (4) SALA.

As shown in Table 4, no single component matches the full SALA objective. Freq. Only confirms the value of prior correction but leaves semantic crowdedness unresolved. Geom. Only improves over the baseline, by demonstrating that  $\Omega_c$  carries independent structural signal, yet lacks the frequency calibration needed for tail-vs-head separation. Their combination yields consistent additive gains that widen with imbalance severity by confirming that the two signals are complementary.

## 5. Conclusion

In this work, we propose Semantic-Aware Logit Adjustment (SALA) that augments the standard frequency-based correction with a semantic crowdedness penalty  $\Omega_c$  computed directly from the cosine similarity between classifier prototypes. By jointly accounting for class rarity and semantic crowdedness, it selectively penalizes classes that are both underrepresented and structurally entangled with their neighbors rather than treating all the tail classes as uniformly difficult. The resulting objective is lightweight and plug-and-play, by requiring no auxiliary networks or contrastive branches.

Experiments on ImageNet-LT, CIFAR-10-LT, CIFAR-100-LT, and iNaturalist 2018 show that SALA consistently outperforms strong frequency-based, margin-based, and contrastive baselines, with especially strong gains on the more fine-grained datasets.

## References

Shaden Alshammari, Yu-Xiong Wang, Deva Ramanan, and Shu Kong. Long-tailed recognition via weight balancing. 03 2022. doi: 10.48550/arXiv.2203.14197.

# Proceedings Track

- Kaidi Cao, Colin Wei, Adrien Gaidon, Nikos Arechiga, and Tengyu Ma. Learning imbalanced datasets with label-distribution-aware margin loss, 2019. URL <https://arxiv.org/abs/1906.07413>.
- N. V. Chawla, K. W. Bowyer, L. O. Hall, and W. P. Kegelmeyer. Smote: Synthetic minority over-sampling technique. *Journal of Artificial Intelligence Research*, 16:321–357, 2002. ISSN 1076-9757. doi: 10.1613/jair.953. URL <http://dx.doi.org/10.1613/jair.953>.
- M. Chen, Y. Du, W. Jiang, B. Zhang, S. Feng, Y. Xin, and C. Wang. Robust logit adjustment for learning with long-tailed noisy data. In *Proceedings of the AAAI Conference on Artificial Intelligence*, volume 39, pages 15830–15838, 2025. doi: 10.1609/aaai.v39i15.33738. URL <https://doi.org/10.1609/aaai.v39i15.33738>.
- Qiong Chen, Tianlin Huang, and Qingfa Liu. Swrm: Similarity window reweighting and margin for long-tailed recognition. *ACM Transactions on Multimedia Computing, Communications, and Applications*, 20, 02 2024. doi: 10.1145/3643816.
- Jang Hyun Cho and Philipp Krähenbühl. Long-tail detection with effective class-margins, 2023. URL <https://arxiv.org/abs/2301.09724>.
- Yin Cui, Menglin Jia, Tsung-Yi Lin, Yang Song, and Serge Belongie. Class-balanced loss based on effective number of samples. In *Proceedings of the IEEE Conference on Computer Vision and Pattern Recognition (CVPR)*, 2019.
- Zongyong Deng, Hao Liu, Yaoxing Wang, Chenyang Wang, Zekuan Yu, and Xuehong Sun. Pml: Progressive margin loss for long-tailed age classification, 2021. URL <https://arxiv.org/abs/2103.02140>.
- Georgios Douzas and Fernando Baao. Self-organizing map oversampling (somo) for imbalanced data set learning. *Expert Systems with Applications*, 82:40–52, 10 2017. doi: 10.1016/j.eswa.2017.03.073.
- Chaoqun Du, Yulin Wang, Shiji Song, and Gao Huang. Probabilistic contrastive learning for long-tailed visual recognition, 2024. URL <https://arxiv.org/abs/2403.06726>.
- Hongyu Guo and Herna Lydia Viktor. Learning from imbalanced data sets with boosting and data generation: the databoost-im approach. *SIGKDD Explor.*, 6:30–39, 2004. URL <https://api.semanticscholar.org/CorpusID:18057117>.
- Haibo He, Yang Bai, Eduardo A. Garcia, and Shutao Li. Adasyn: Adaptive synthetic sampling approach for imbalanced learning. In *2008 IEEE International Joint Conference on Neural Networks (IEEE World Congress on Computational Intelligence)*, pages 1322–1328, 2008. doi: 10.1109/IJCNN.2008.4633969.
- Youngkyu Hong, Seungju Han, Kwanghee Choi, Seokjun Seo, Beomsu Kim, and Buru Chang. Disentangling label distribution for long-tailed visual recognition, 2021. URL <https://arxiv.org/abs/2012.00321>.

# Proceedings Track

- Grant Van Horn, Oisín Mac Aodha, Yang Song, Yin Cui, Chen Sun, Alex Shepard, Hartwig Adam, Pietro Perona, and Serge Belongie. The inaturalist species classification and detection dataset, 2018. URL <https://arxiv.org/abs/1707.06642>.
- Chengkai Hou, Jieyu Zhang, Haonan Wang, and Tianyi Zhou. Subclass-balancing contrastive learning for long-tailed recognition, 2023. URL <https://arxiv.org/abs/2306.15925>.
- Cui Jiequan, Shu Liu, Zhuotao Tian, Zhisheng Zhong, and Jiaya Jia. Reslt: Residual learning for long-tailed recognition. *IEEE Transactions on Pattern Analysis and Machine Intelligence*, PP:1–1, 05 2022. doi: 10.1109/TPAMI.2022.3174892.
- Amelia Jiménez-Sánchez, Natalia-Rozalia Avlona, Sarah de Boer, Víctor M. Campello, Aasa Feragen, Enzo Ferrante, Melanie Ganz, Judy Wawira Gichoya, Camila Gonzalez, Steff Groefsema, Alessa Hering, Adam Hulman, Leo Joskowicz, Dovile Juodelyte, Melih Kandemir, Thijs Kooi, Jorge del Pozo Lérída, Livie Yumeng Li, Andre Pacheco, Tim Rädsch, Mauricio Reyes, Théo Sourget, Bram van Ginneken, David Wen, Nina Weng, Jack Junchi Xu, Hubert Dariusz Zającz, Maria A. Zuluaga, and Veronika Cheplygina. In the picture: Medical imaging datasets, artifacts, and their living review. In *Proceedings of the 2025 ACM Conference on Fairness, Accountability, and Transparency*, FAccT ’25, page 511–531. ACM, June 2025. doi: 10.1145/3715275.3732035. URL <http://dx.doi.org/10.1145/3715275.3732035>.
- Bingyi Kang, Saining Xie, Marcus Rohrbach, Zhicheng Yan, Albert Gordo, Jiashi Feng, and Yannis Kalantidis. Decoupling representation and classifier for long-tailed recognition. In *International Conference on Learning Representations (ICLR)*, 2020.
- Sota Kato and Kazuhiro Hotta. Enlarged large margin loss for imbalanced classification, 2023. URL <https://arxiv.org/abs/2306.09132>.
- Prannay Khosla, Piotr Teterwak, Chen Wang, Aaron Sarna, Yonglong Tian, Phillip Isola, Aaron Maschinot, Ce Liu, and Dilip Krishnan. Supervised contrastive learning, 2021. URL <https://arxiv.org/abs/2004.11362>.
- Alex Krizhevsky. Learning multiple layers of features from tiny images. 2009. URL <https://api.semanticscholar.org/CorpusID:18268744>.
- Mengke Li, Yiu-Ming Cheung, and Yang Lu. Long-tailed visual recognition via gaussian clouded logit adjustment. In *2022 IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*, page 6919–6928. IEEE, 2022a. doi: 10.1109/cvpr52688.2022.00680. URL <http://dx.doi.org/10.1109/CVPR52688.2022.00680>.
- Mengke Li, Yiu-ming Cheung, Yang Lu, Zhikai Hu, Weichao Lan, and Hui Huang. Adjusting logit in gaussian form for long-tailed visual recognition. *IEEE Transactions on Artificial Intelligence*, 5(10):5026–5039, October 2024. ISSN 2691-4581. doi: 10.1109/tai.2024.3401102. URL <http://dx.doi.org/10.1109/TAI.2024.3401102>.
- Shuang Li, Kaixiong Gong, Chi Harold Liu, Yulin Wang, Feng Qiao, and Xinjing Cheng. Metasaug: Meta semantic augmentation for long-tailed visual recognition, 2021. URL <https://arxiv.org/abs/2103.12579>.

# Proceedings Track

- Tianhong Li, Peng Cao, Yuan Yuan, Lijie Fan, Yuzhe Yang, Rogerio Feris, Piotr Indyk, and Dina Katabi. Targeted supervised contrastive learning for long-tailed recognition, 2022b. URL <https://arxiv.org/abs/2111.13998>.
- Tsung-Yi Lin, Priya Goyal, Ross Girshick, Kaiming He, and Piotr Dollár. Focal loss for dense object detection, 2018. URL <https://arxiv.org/abs/1708.02002>.
- Mingyu Liu, Ekim Yurtsever, Jonathan Fossaert, Xingcheng Zhou, Walter Zimmer, Yuning Cui, Bare Luka Zagar, and Alois C. Knoll. A survey on autonomous driving datasets: Statistics, annotation quality, and a future outlook, 2024. URL <https://arxiv.org/abs/2401.01454>.
- Ziwei Liu, Zhongqi Miao, Xiaohang Zhan, Jingdong Wang, Boqing Gong, and Stella X. Yu. Large-scale long-tailed recognition in an open world. In *Proceedings of the IEEE Conference on Computer Vision and Pattern Recognition (CVPR)*, 2019.
- Aditya Krishna Menon, Sadeep Jayasumana, Ankit Singh Rawat, Himanshu Jain, Andreas Veit, and Sanjiv Kumar. Long-tail learning via logit adjustment. In *International Conference on Learning Representations (ICLR)*, 2021.
- Muhammad Mujahid, Erol Kina, Furqan Rustam, Mònica Gracia, Eduardo Alvarado, Isabel De la Torre Díez, and Imran Ashraf. Data oversampling and imbalanced datasets: an investigation of performance for machine learning and feature engineering. *Journal of Big Data*, 11, 06 2024. doi: 10.1186/s40537-024-00943-4.
- Vardan Papayan, X. Y. Han, and David L. Donoho. Prevalence of neural collapse during the terminal phase of deep learning training. *Proceedings of the National Academy of Sciences*, 117(40):24652–24663, 2020. ISSN 1091-6490. doi: 10.1073/pnas.2015509117. URL <http://dx.doi.org/10.1073/pnas.2015509117>.
- Seulki Park, Jongin Lim, Younghun Jeon, and Jin Young Choi. Influence-balanced loss for imbalanced visual classification, 2021. URL <https://arxiv.org/abs/2110.02444>.
- Kunpeng Qiu, Zhiying Zhou, and Yongxin Guo. Learn from zoom: Decoupled supervised contrastive learning for wce image classification, 2024. URL <https://arxiv.org/abs/2401.05771>.
- Jiawei Ren, Bingbing Yu, Meng Wang, and Xiaolin Zhao. Balanced softmax for long-tailed visual recognition. In *Advances in Neural Information Processing Systems (NeurIPS)*, 2020.
- Olga Russakovsky, Jia Deng, Hao Su, Jonathan Krause, Sanjeev Satheesh, Sean Ma, Zhiheng Huang, Andrej Karpathy, Aditya Khosla, Michael Bernstein, Alexander C. Berg, and Li Fei-Fei. Imagenet large scale visual recognition challenge, 2015. URL <https://arxiv.org/abs/1409.0575>.
- Jürgen Schmidhuber. Deep learning in neural networks: An overview. *Neural Networks*, 61:85–117, January 2015. ISSN 0893-6080. doi: 10.1016/j.neunet.2014.09.003. URL <http://dx.doi.org/10.1016/j.neunet.2014.09.003>.

# Proceedings Track

- Minseok Son, Inyong Koo, Jinyoung Park, and Changick Kim. Difficulty-aware balancing margin loss for long-tailed recognition, 2024. URL <https://arxiv.org/abs/2412.15477>.
- Jingyuan Tan, Changhong Wang, Bin Liu, Qi Wu, and Yandong Guo. Equalization loss for long-tailed object recognition. In *Proceedings of the IEEE Conference on Computer Vision and Pattern Recognition (CVPR)*, 2020.
- Yingfan Tao, Jingna Sun, Hao Yang, Li Chen, Xu Wang, Wenming Yang, Daniel Du, and Min Zheng. Local and global logit adjustments for long-tailed learning. In *2023 IEEE/CVF International Conference on Computer Vision (ICCV)*, pages 11749–11758, 2023. doi: 10.1109/ICCV51070.2023.01082.
- Elsbeth Turcan and Kathy McKeown. Dreddit: A Reddit dataset for stress analysis in social media. In Eben Holderness, Antonio Jimeno Yepes, Alberto Lavelli, Anne-Lyse Minard, James Pustejovsky, and Fabio Rinaldi, editors, *Proceedings of the Tenth International Workshop on Health Text Mining and Information Analysis (LOUHI 2019)*, pages 97–107, Hong Kong, November 2019. Association for Computational Linguistics. doi: 10.18653/v1/D19-6213. URL <https://aclanthology.org/D19-6213/>.
- Jiaqi Wang, Wenwei Zhang, Yuhang Zang, Yuhang Cao, Jiangmiao Pang, Tao Gong, Kai Chen, Ziwei Liu, Chen Change Loy, and Dahua Lin. Seesaw loss for long-tailed instance segmentation, 2021. URL <https://arxiv.org/abs/2008.10032>.
- Yidong Wang, Bowen Zhang, Wenxin Hou, Zhen Wu, Jindong Wang, and Takahiro Shinozaki. Margin calibration for long-tailed visual recognition, 2022. URL <https://arxiv.org/abs/2112.07225>.
- Liang Xie, Yibo Yang, Deng Cai, and Xiaofei He. Neural collapse inspired attraction–repulsion-balanced loss for imbalanced learning. *Neurocomputing*, 527:60–70, March 2023. ISSN 0925-2312. doi: 10.1016/j.neucom.2023.01.023. URL <http://dx.doi.org/10.1016/j.neucom.2023.01.023>.
- Hongren Yan, Yuhua Qian, Furong Peng, Jiachen Luo, Zheqing Zhu, and Feijiang Li. Neural collapse to multiple centers for imbalanced data. In A. Globerson, L. Mackey, D. Belgrave, A. Fan, U. Paquet, J. Tomczak, and C. Zhang, editors, *Advances in Neural Information Processing Systems*, volume 37, pages 65583–65617. Curran Associates, Inc., 2024. doi: 10.52202/079017-2095. URL [https://proceedings.neurips.cc/paper\\_files/paper/2024/file/78a2ed8e7da1829d64fb3561adbd706b-Paper-Conference.pdf](https://proceedings.neurips.cc/paper_files/paper/2024/file/78a2ed8e7da1829d64fb3561adbd706b-Paper-Conference.pdf).
- Yibo Yang, Shixiang Chen, Xiangtai Li, Liang Xie, Zhouchen Lin, and Dacheng Tao. Inducing neural collapse in imbalanced learning: Do we really need a learnable classifier at the end of deep neural network?, 2022. URL <https://arxiv.org/abs/2203.09081>.
- Sihao Yu, Jiafeng Guo, Ruqing Zhang, Yixing Fan, Zizhen Wang, and Xueqi Cheng. A re-balancing strategy for class-imbalanced classification based on instance difficulty. pages 70–79, 06 2022. doi: 10.1109/CVPR52688.2022.00017.

# Proceedings Track

- Enhao Zhang, Chuanxing Geng, Chaohua Li, and Songcan Chen. Dynamic learnable logit adjustment for long-tailed visual recognition. *IEEE Transactions on Circuits and Systems for Video Technology*, 34(9):7986–7997, 2024a. doi: 10.1109/TCSVT.2024.3383962.
- Fan Zhang, Wei Qin, Weijieying Ren, Lei Wang, Zetong Chen, and Richang Hong. Gradient-aware logit adjustment loss for long-tailed classifier, 2024b. URL <https://arxiv.org/abs/2403.09036>.
- Shaoyu Zhang, Chen Chen, Xiyuan Hu, and Silong Peng. Balanced knowledge distillation for long-tailed learning. *Neurocomputing*, 527, 03 2023. doi: 10.1016/j.neucom.2023.01.063.
- Yan Zhao, Weicong Chen, Xu Tan, Kai Huang, and Jihong Zhu. Adaptive logit adjustment loss for long-tailed visual recognition, 2021. URL <https://arxiv.org/abs/2104.06094>.
- Zhisheng Zhong, Cui Jiequan, Shu Liu, and Jiaya Jia. Improving calibration for long-tailed recognition. pages 16484–16493, 06 2021. doi: 10.1109/CVPR46437.2021.01622.
- Xiaoling Zhou, Ou Wu, and Nan Yang. Class and attribute-aware logit adjustment for generalized long-tail learning. In *Proceedings of the Thirty-Ninth AAAI Conference on Artificial Intelligence and Thirty-Seventh Conference on Innovative Applications of Artificial Intelligence and Fifteenth Symposium on Educational Advances in Artificial Intelligence*, AAAI’25/IAAI’25/EAAI’25. AAAI Press, 2025. ISBN 978-1-57735-897-8. doi: 10.1609/aaai.v39i21.34462. URL <https://doi.org/10.1609/aaai.v39i21.34462>.
- Jianggang Zhu, Zheng Wang, Jingjing Chen, Yi-Ping Phoebe Chen, and Yu-Gang Jiang. Balanced contrastive learning for long-tailed visual recognition, 2022. URL <https://arxiv.org/abs/2207.09052>.

# Proceedings Track

## Appendix A. Computational Overhead Analysis

We analyze the computational cost of SALA relative to standard Cross-Entropy (CE) and Logit Adjustment (LA). SALA introduces a semantic penalty based on pairwise cosine similarities between classifier prototypes. Let  $C$  denote the number of classes,  $d$  the feature dimension, and  $B$  the mini-batch size.

**Complexity.** The standard CE loss requires a backbone forward pass followed by a linear classifier with a complexity  $O(B \cdot C \cdot d)$ . SALA adds two classifier-head operations: (i)  $L_2$ -normalization of the classifier weights, which costs  $O(C \cdot d)$ , and (ii) construction of the prototype similarity matrix  $S = \bar{W}\bar{W}^\top \in \mathbb{R}^{C \times C}$ , which costs  $O(C^2 \cdot d)$  when recomputed from scratch. In practice, these operations act only on the classifier head and are therefore much cheaper than the convolutional backbone.

**Dataset-specific implementation.** For CIFAR-10-LT, CIFAR-100-LT, and ImageNet-LT, the prototype similarity matrix is refreshed every iteration. For iNaturalist2018, where  $C = 8,142$ , we reduce the cost by **refreshing the prototype similarity matrix once every 100 iterations** and reusing the cached semantic scores between refreshes. This makes the expensive  $O(C^2 \cdot d)$  computation amortized over multiple updates and prevents the semantic term from becoming a bottleneck on the largest benchmark.

Table 5: Computational overhead of SALA relative to the backbone forward pass. Overhead is reported per training iteration, with the iNaturalist2018 prototype similarity matrix refreshed every 100 iterations.

| Dataset       | Backbone  | Classes ( $C$ ) | Dim ( $d$ ) | Refresh Rate    | Backbone FLOPs   | SALA Overhead                |
|---------------|-----------|-----------------|-------------|-----------------|------------------|------------------------------|
| CIFAR-10-LT   | ResNet-32 | 10              | 64          | every iter.     | $\approx 0.07$ G | $< 0.01\%$                   |
| CIFAR-100-LT  | ResNet-32 | 100             | 64          | every iter.     | $\approx 0.07$ G | $\approx 0.08\%$             |
| ImageNet-LT   | ResNet-50 | 1,000           | 2,048       | every iter.     | $\approx 3.80$ G | $\approx 0.24\%$             |
| iNaturalist18 | ResNet-50 | 8,142           | 2,048       | every 100 iters | $\approx 3.80$ G | $\approx 0.14\%$ (amortized) |

**Scalability.** For extreme classification tasks where  $C$  may reach tens of thousands, re-computing the full similarity matrix at every step can become really expensive. SALA therefore incorporates periodic prototype updates, i.e., recomputing  $S$  once every  $K$  steps or once per epoch. This reduces the amortized similarity cost from  $O(C^2 \cdot d)$  per iteration to  $O(C^2 \cdot d/K)$ , while preserving the geometric effect of the semantic bias.

**Empirical observations.** In our experiments on an NVIDIA A100 GPU, we observed no measurable difference in wall-clock training time per epoch between the baseline LA and SALA ( $\pm 1.2$  seconds for ImageNet-LT). This indicates that SALA provides the intended geometric benefits without materially increasing the training budget.

## Appendix B. Training Procedure

We use two different optimization techniques to handle the semantic penalty and maintain the training stability. At first during the forward pass, we detach the classifier weights from the computational graph; this it fixes  $\Omega_c$  as a constant scalar for that particular

# Proceedings Track

iteration. Making sure that  $\nabla_W \Omega_c = 0$ . Therefore, gradients only flow through the logit representations to update the network parameters and prevent chaotic optimization loops. Second, we follow a two-stage training process where initially the model runs with  $\alpha = 0$ , which lets class clusters form based on simple frequency counts. When these initial spaces become stable, we then turn on the SALA penalty to untangle the fine-grained geometric issues affecting the tail classes.